



March 2015

Observable Applications w/ Node.js

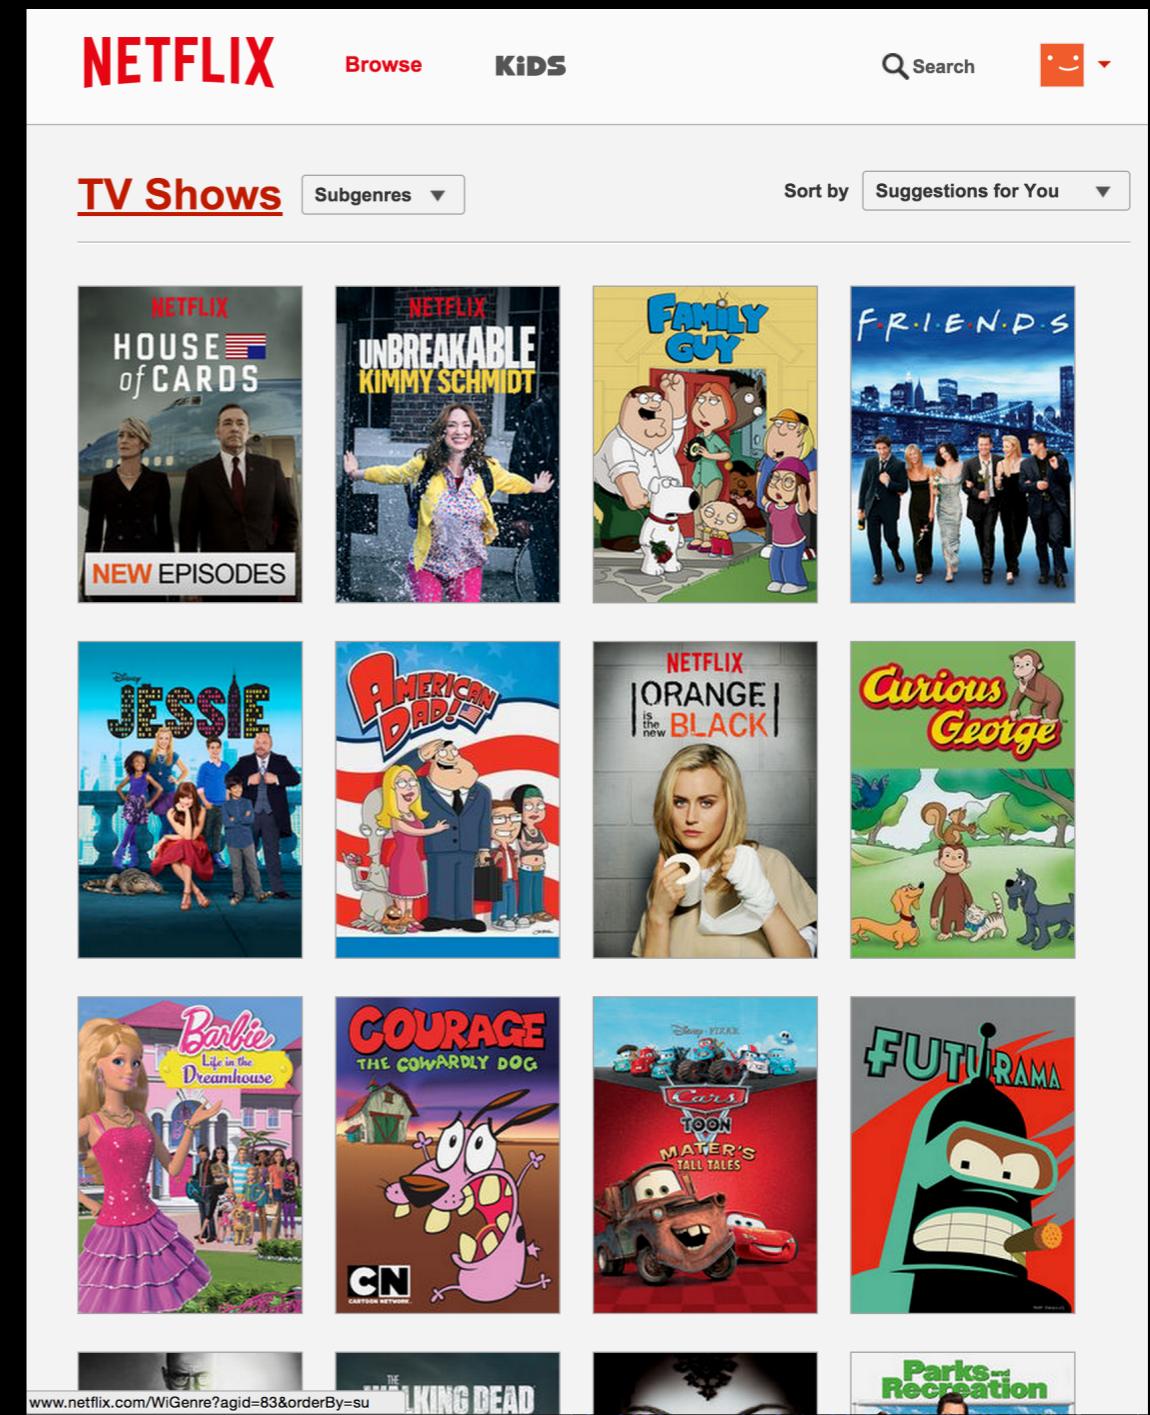
Yunong Xiao, Senior Node.js Engineer,
UI Platform, Netflix
[@yunongx, yunong@netflix.com](mailto:yunong@netflix.com)



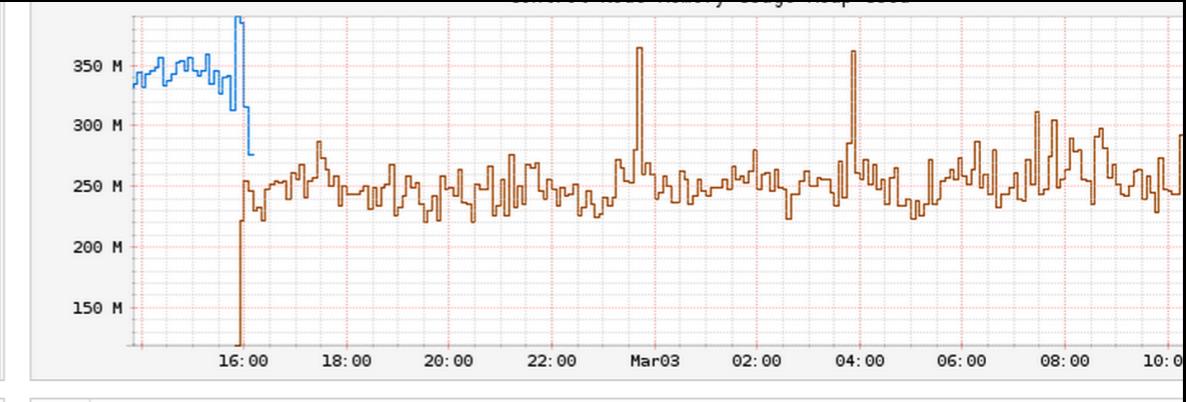
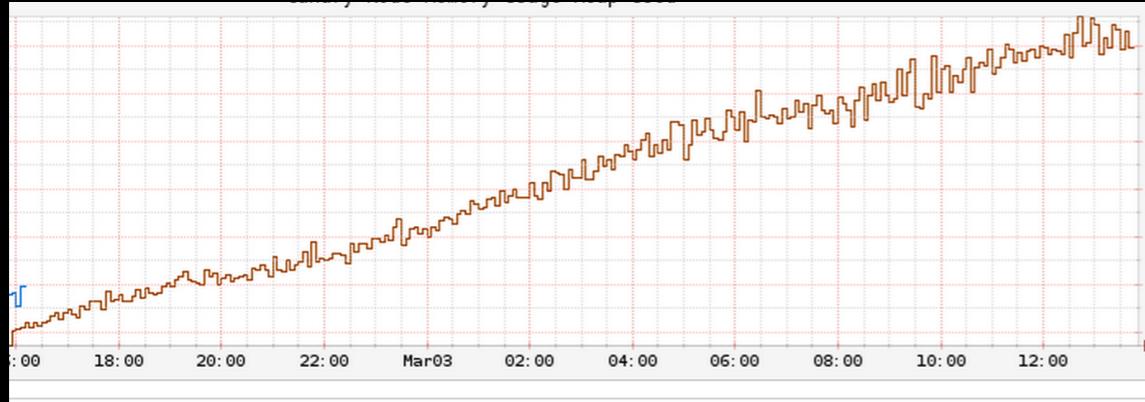
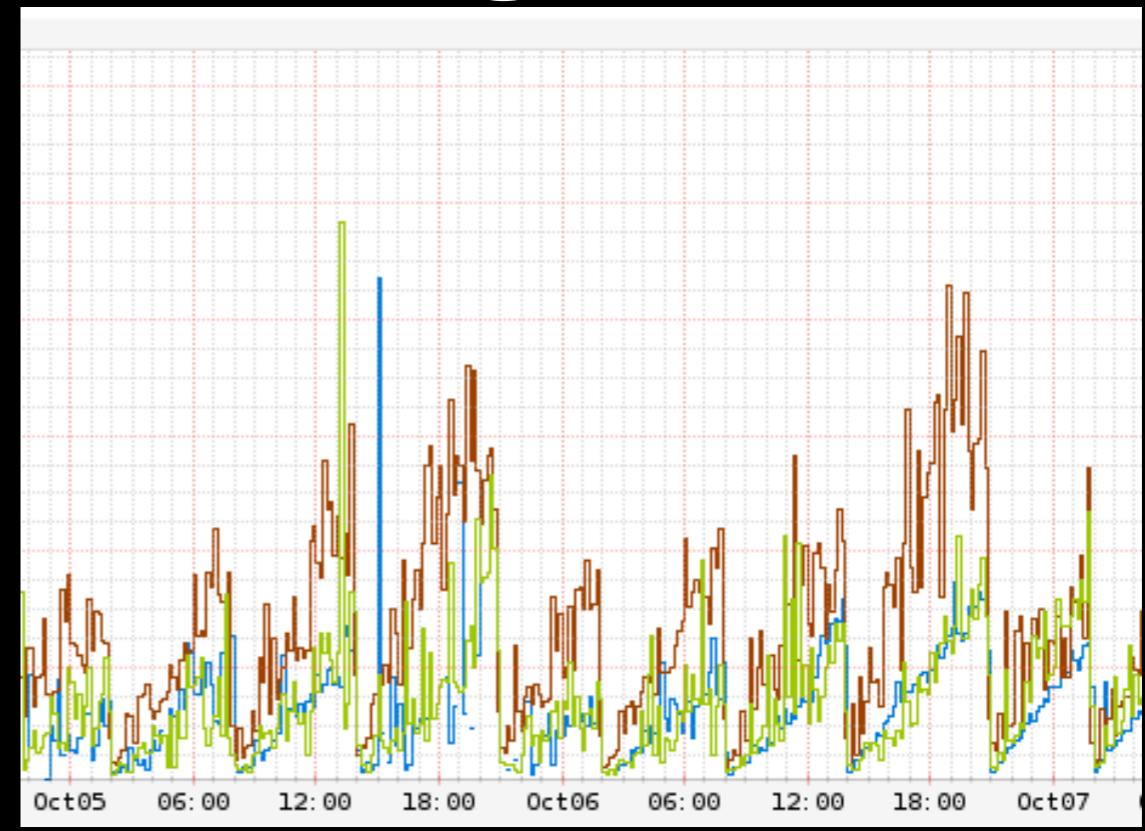
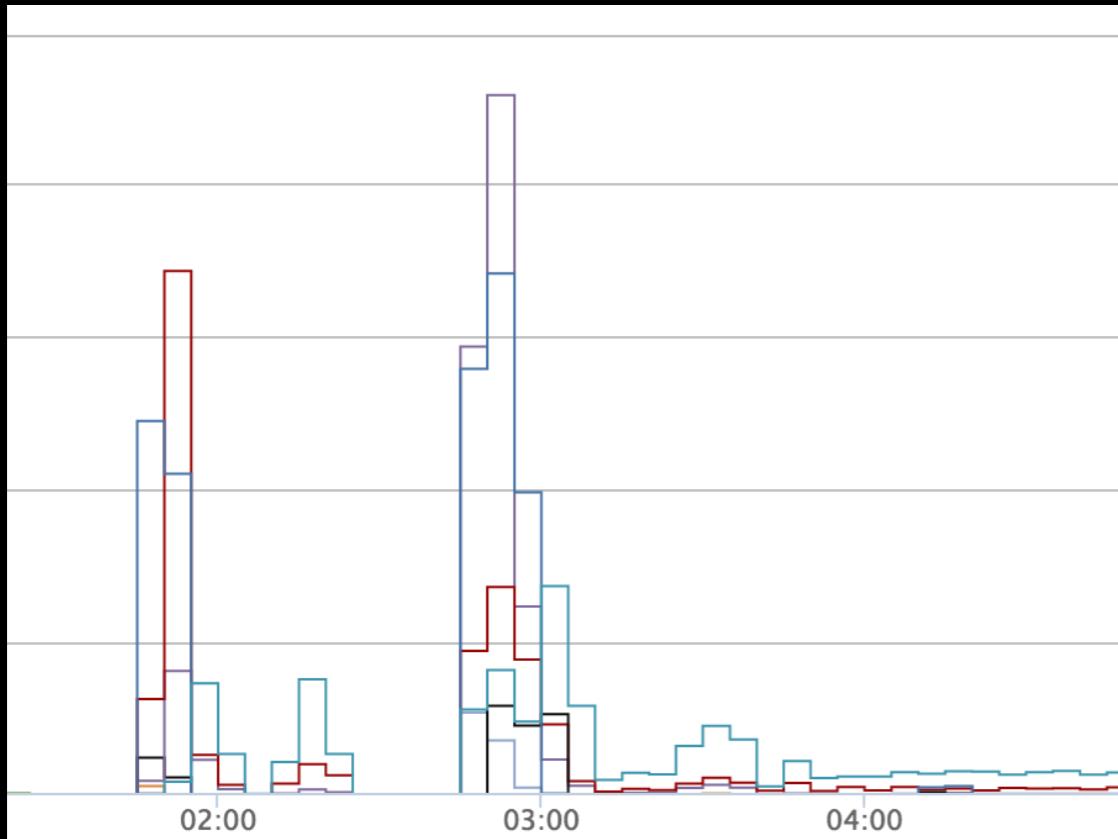
NETFLIX

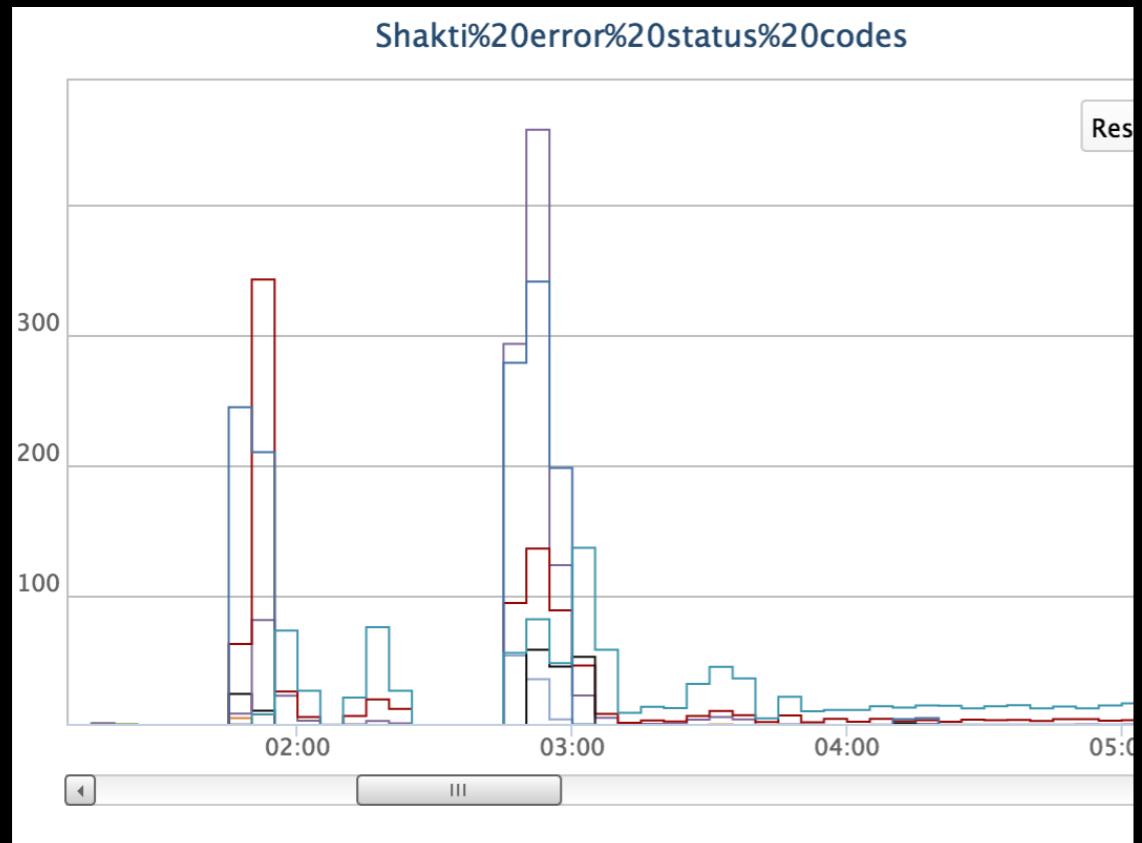
Node.js @ Netflix

- Website www.netflix.com
 - Discovery
 - Playback
 - Acquisition
 - Account Management
- Internal Services
- TV/Mobile/Devices

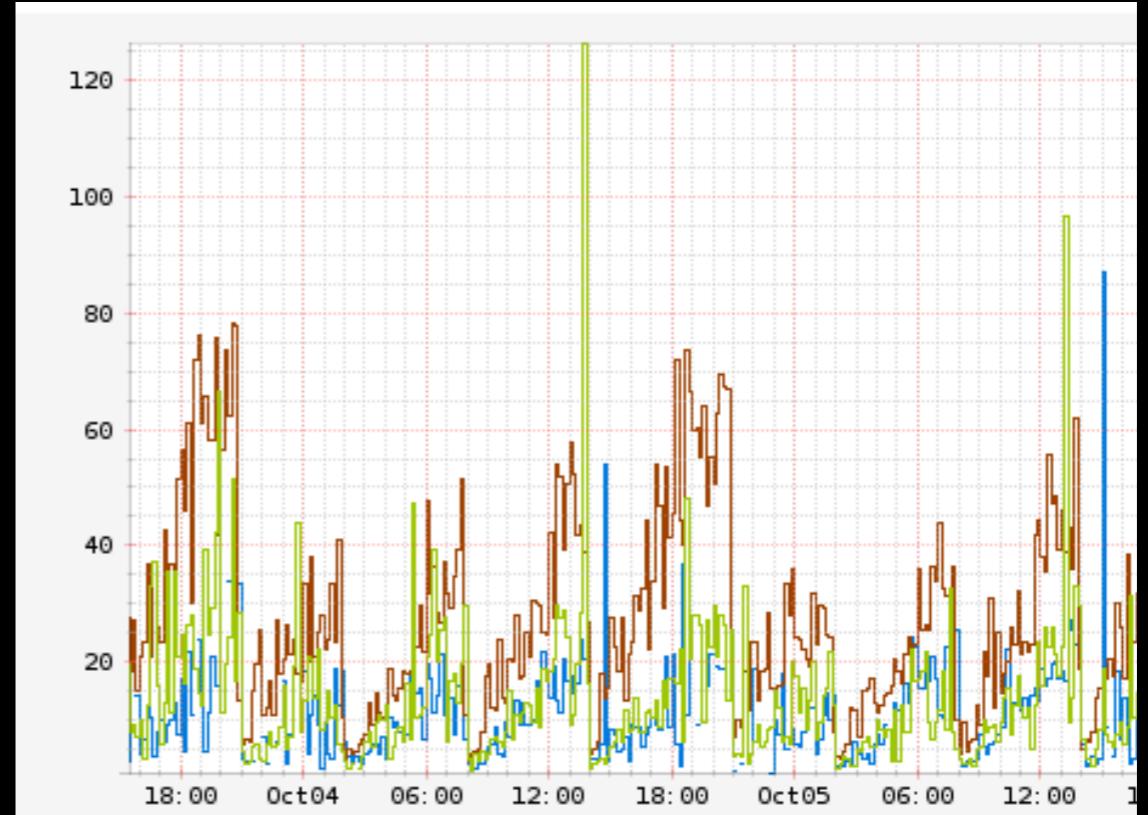


What's Wrong?

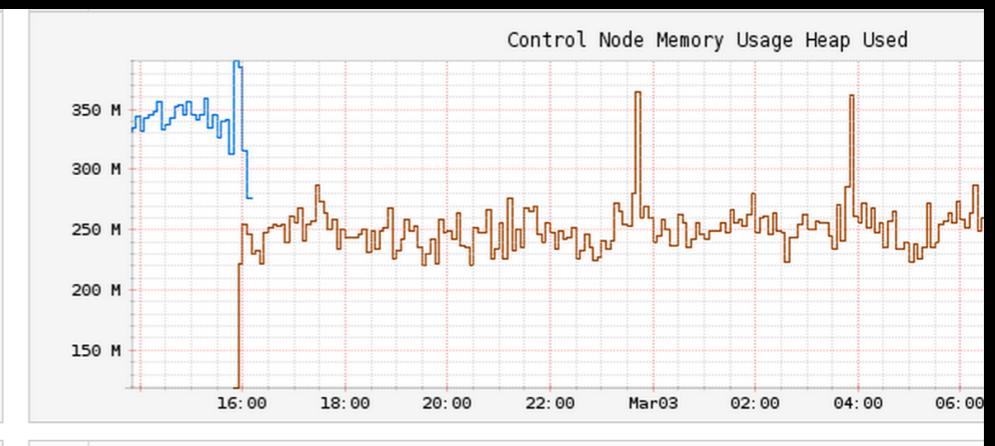
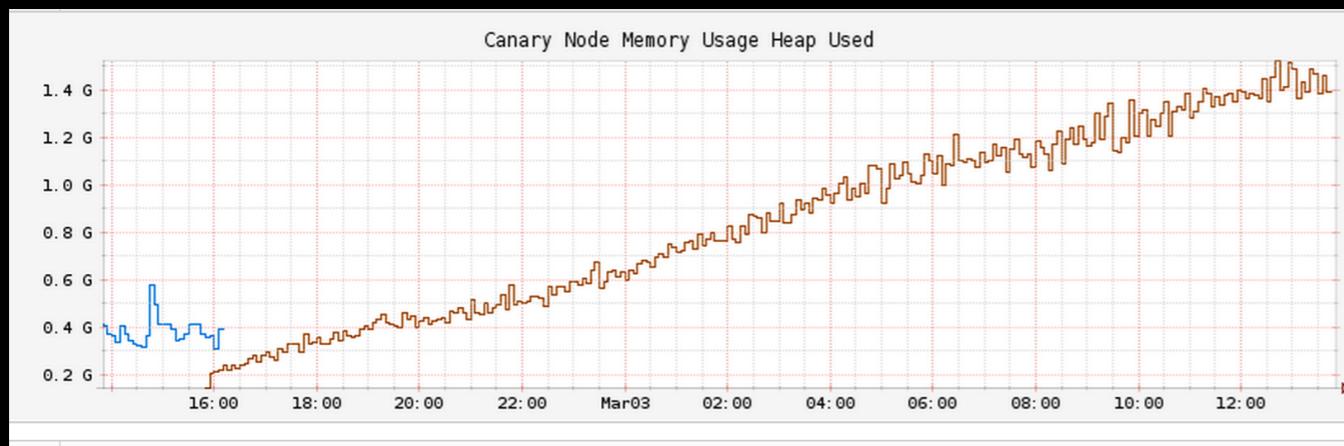




Increased Errors



Increased Latency



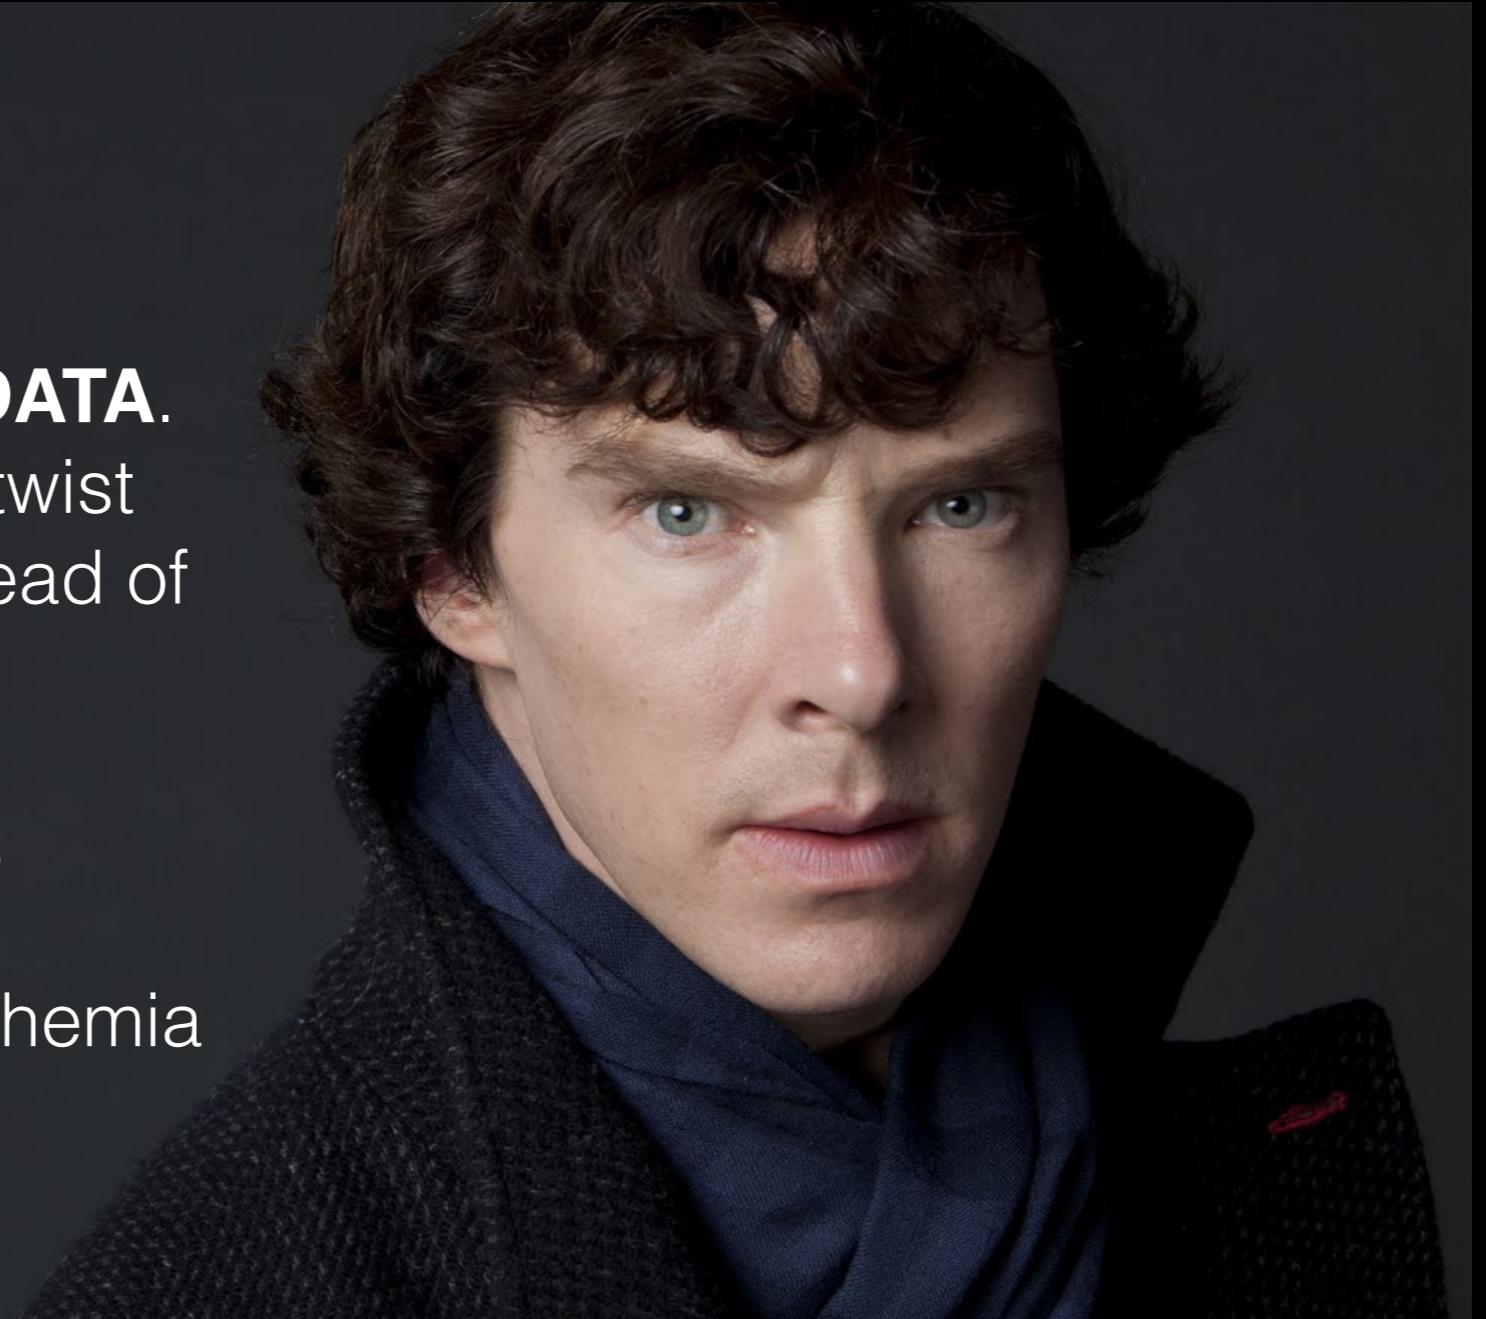
Memory Leak

What Now?

“It is a capital mistake to theorize before one has **DATA**. Insensibly one begins to twist facts to suit theories, instead of theories to suit facts.”

Sherlock Holmes

-A Scandal in Bohemia



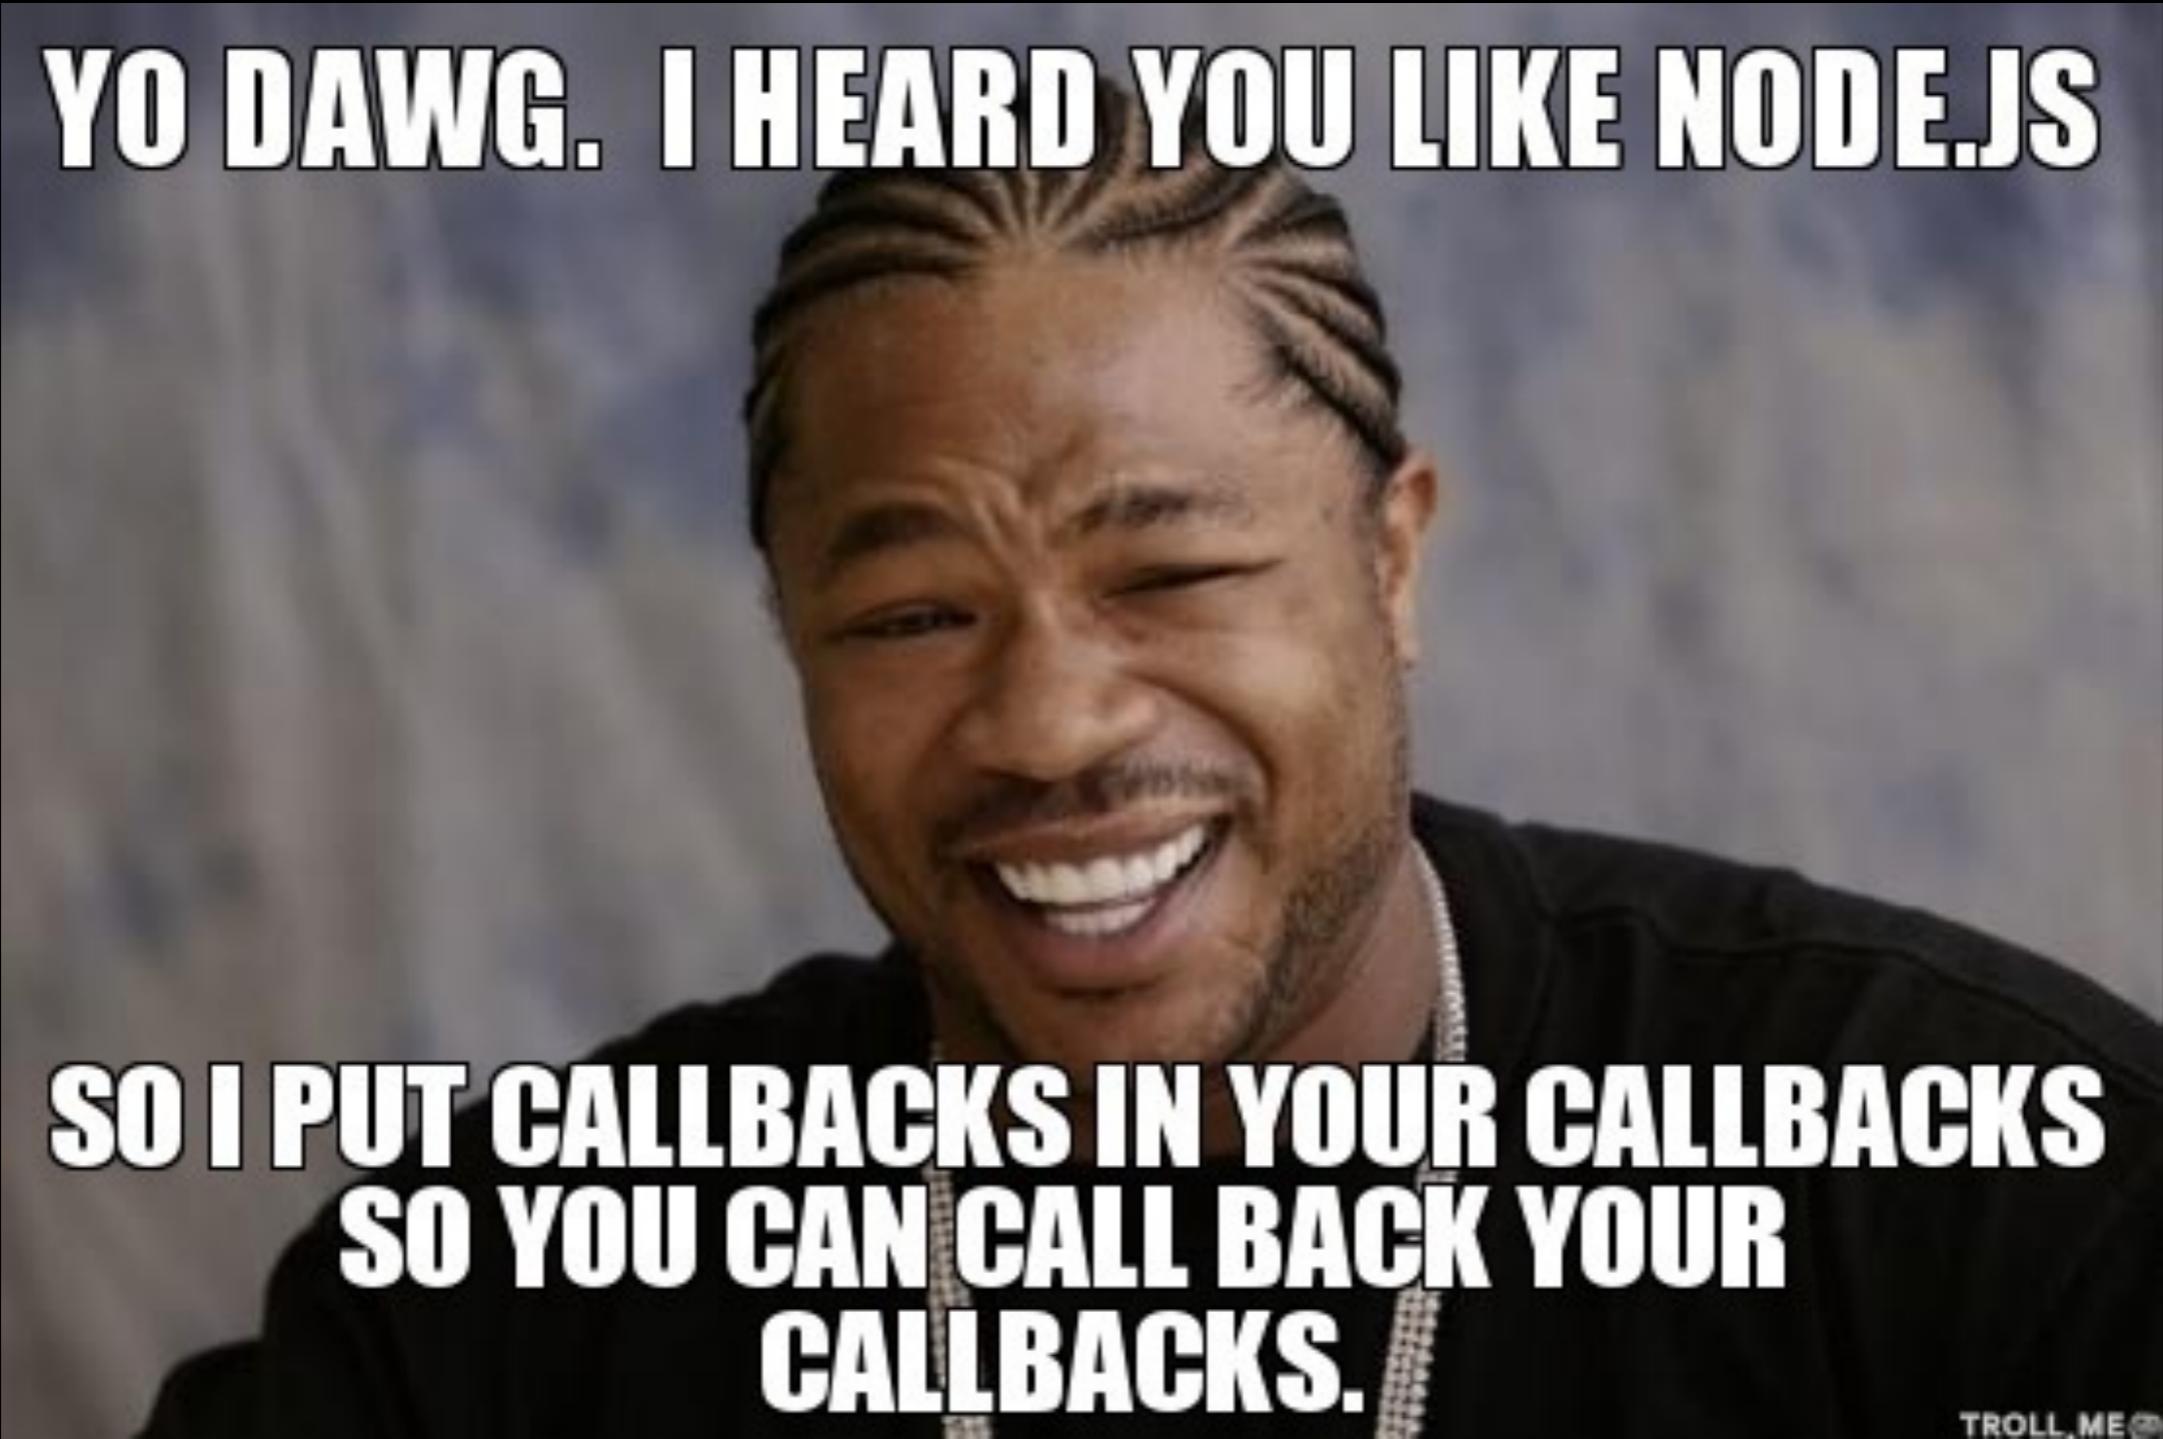
Data Data Data





Node.js Modules

```
1 var db = require('./db/callbackDb');
2
3 db.set('key1', 'value1', function(err) {
4     if(err) throw err;
5
6     db.set('key2', 'value2', function(err) {
7         if(err) throw err;
8
9         db.set('key3', 'value3', function(err) {
10            if(err) throw err;
11
12            var str = '';
13            db.get('key1', function(err, value) {
14                if(err) throw err;
15
16                str += value + ' - ';
17
18                console.log(str);
19            });
20        });
21    });
22});
```



YO DAWG. I HEARD YOU LIKE NODE.JS

**SO I PUT CALLBACKS IN YOUR CALLBACKS
SO YOU CAN CALL BACK YOUR
CALLBACKS.**

```
1 var vasync = require('vasync');
2
3 var results = vasync.pipeline({funcs: [
4     function one(_, cb) {
5         // some async function one
6         return cb();
7     },
8     function two(_, cb) {
9         // some async function two
10        return cb();
11    },
12    function three(_, cb) {
13        // some async function three
14        return cb();
15    }
16 ], args: {}}, function cb(err, res) {
17     console.log('finished pipeline', err, res);
18});
```

vasync

- <https://github.com/davepacheco/node-vasync>
- Similar to caolan/async.
- Added observability.

```
async f() state  
{ operations:  
  [ { func: [Function: one],  
    funcname: 'one',  
    status: 'fail', ← failed f()  
    err: [Error],  
    result: undefined },  
   { func: [Function: two],  
     funcname: 'two',  
     status: 'ok', ← successful f()  
     err: undefined,  
     result: undefined },  
   { func: [Function: three], funcname: 'three', status: 'pending' } ],  
  successes: [ undefined ],  
  ndone: 2, ← finished f()  
  nerrors: 1 }  
# of errors
```

vasync

Async workflow management is error prone

- Make API requests.
- Persist state to file system.
- Query database.
- Update caches.

Example

```
1 var vasync = require('vasync');
2
3 var results = vasync.pipeline({funcs: [
4     function one(_, cb) {
5         console.log('some async function one');
6         return cb();
7     },
8     function two(_, cb) {
9         console.log('some async function two');
10        return cb();
11    },
12    function three(_, cb) {
13        console.log('some async function three');
14        // No cb due to bug
15        // return cb(); ← cb() not invoked
16    }
17 ], args: {}}, function cb(err, res) {
18     console.log('finished pipeline', err, res);
19 });
20
21 setInterval(function() { console.log('pipeline status', results); }, 1000);
22
```

Results

```
{ operations:  
  [ { func: [Function: one],  
      funcname: 'one',  
      status: 'ok',  
      err: undefined,  
      result: undefined },  
    { func: [Function: two],  
      funcname: 'two',  
      status: 'ok',  
      err: undefined,  
      result: undefined },  
    { func: [Function: three], funcname: 'three', status: 'pending' } ],  
  successes: [ undefined, undefined ],  
  ndone: 2,  
  nerrors: 0 }
```

function three pending



How Do I See this in Prod?

- Logs
- Core Dumps
- REPL
- HTTP API

node-bunyan

- <https://github.com/trentm/node-bunyan>
- Streaming JSON logging library for JS. (node and browser)

Streaming JSON

- One JSON object per line. e.g. Twitter's streaming API.
- Perfect for machine processing. i.e. works with Unix tools such as grep(1), cut(1), ...
- Illegible for Humans

Streaming JSON

```
{"name": "shakti", "hostname": "shakti-prod-i-232a3529", "pid": 7567, "component": "restify", "req_id": "9d92d710-b961-11e4-a9bd-a188e49c3b0c", "level": 30, "method": "GET", "fullResourceUrl": "/shakti/90409bc7/memberContext", "statusCode": 200, "elapsedTime": 305, "msg": "[nf-httplib]: memberContext request success", "time": "2015-02-21T00:36:03.763Z", "v": 0} {"name": "shakti", "hostname": "shakti-prod-i-232a3529", "pid": 7567, "component": "restify", "req_id": "9d92d710-b961-11e4-a9bd-a188e49c3b0c", "level": 30, "msg": "Resolving action shard: search", "time": "2015-02-21T00:36:03.763Z", "v": 0} {"name": "shakti", "hostname": "shakti-prod-i-232a3529", "pid": 7567, "component": "restify", "req_id": "9d8f0680-b961-11e4-a9bd-a188e49c3b0c", "level": 30, "method": "GET", "fullResourceUrl": "/shakti/90409bc7/memberContext", "statusCode": 200, "elapsedTime": 325, "msg": "[nf-httplib]: memberContext request success", "time": "2015-02-21T00:36:03.767Z", "v": 0} {"name": "shakti", "hostname": "shakti-prod-i-232a3529", "pid": 7567, "component": "restify", "req_id": "9d3aa4f0-b961-11e4-a9bd-a188e49c3b0c", "level": 30, "method": "GET", "fullResourceUrl": "/shakti/90409bc7/usermessages?authURL=1424478962983.IuqDa5%2BZsu5Iq%2FF4oaWe%2BNJya0s%3D", "statusCode": 200, "elapsedTime": 360, "msg": "[nf-httplib]: umaModel request success", "time": "2015-02-21T00:36:03.769Z", "v": 0}
```

bunyan(1) CLI

```
└─[0] < cat /tmp/fool bunyan
[2015-02-21T00:36:03.763Z] INFO: shakti/restify/7567 on shakti-prod-i-232a3529: [nf-h
ttp]: memberContext request success (req_id=9d92d710-b961-11e4-a9bd-a188e49c3b0c, meth
od=GET, fullResourceUrl=/shakti/90409bc7/memberContext, statusCode=200, elapsedT
ime=30
5)
[2015-02-21T00:36:03.763Z] DEBUG: shakti/restify/7567 on shakti-prod-i-232a3529: Resol
ving action shard: search (req_id=9d92d710-b961-11e4-a9bd-a188e49c3b0c)
[2015-02-21T00:36:03.767Z] TRACE: shakti/restify/7567 on shakti-prod-i-232a3529: [nf-h
ttp]: memberContext request success (req_id=9d8f0680-b961-11e4-a9bd-a188e49c3b0c, meth
od=GET, fullResourceUrl=/shakti/90409bc7/memberContext, statusCode=200, elapsedT
ime=32
5)
[2015-02-21T00:36:03.769Z] WARN: shakti/restify/7567 on shakti-prod-i-232a3529: [nf-h
ttp]: umaModel request success (req_id=9d3aa4f0-b961-11e4-a9bd-a188e49c3b0c, method=GE
T, statusCode=200, elapsedT
ime=360)
    fullResourceUrl: /shakti/90409bc7/usermessages?authURL=1424478962983.IuqDa5%2BZsu5
Iq%2FF4oaWe%2BNJya0s%3D
[2015-02-21T00:36:03.769Z] ERROR: shakti/restify/7567 on shakti-prod-i-232a3529: [nf-h
ttp]: umaModel request success (req_id=9d3aa4f0-b961-11e4-a9bd-a188e49c3b0c, method=GE
T, statusCode=200, elapsedT
ime=360)
    fullResourceUrl: /shakti/90409bc7/usermessages?authURL=1424478962983.IuqDa5%2BZsu5
Iq%2FF4oaWe%2BNJya0s%3D
```

Features

- Lightweight API
- Log levels: trace, debug, info, warn, error, fatal
- Extensible Streams interface.
- Custom object rendering with serializers.
- DTrace support.

Demo

DTrace

- Dynamic tracing framework.
- Available on Mac OS X (dev env). `man dtrace`
- Node.js DTrace USDT provider: <https://github.com/chrisa/node-dtrace-provider>
- Pure JS probes.

bunyan -p
Demo

Production Tracing

- Production environment is Ubuntu Linux on EC2.
- Future work: get probe access. Possibilities:
 - SystemTap
 - LTTng (Linux Tracing Toolkit Next Gen)
 - perf_events markers (currently a proposed patch)

restify

- <http://restifyjs.com>
- Production tested REST framework.
- Observability.
- Metrics.
- First class Bunyan integration.
- DTrace.

Vanilla App

```
1 var restify = require('restify');
2
3 var server = restify.createServer({}); 
4
5 server.use(restify.queryParser());
6
7 server.get(
8     '/hello',
9     function hello(req, res, next) {
10         res.send({hi: req.query.name});
11         return next();
12     }
13 );
14
15 server.listen(1337, function () {
16     console.log('server started', server.url);
17 });
```

Vanilla App

```
└─[1] ◇ node vanilla.js&
[1] 5683
└─[yunong@lgml-yunong] - [~/workspace/restify-demo] - [2015-03-17 12:25:56]
└─[0] ◇ server started

└─[yunong@lgml-yunong] - [~/workspace/restify-demo] - [2015-03-17 12:25:58]
└─[0] ◇ curl 'localhost:1337/hello?name=yunong'
{"hi":"yunong"}%
```

Logging

Request Capture Stream

- Tradeoffs, Perf vs Verbosity.
- Captures all log statements including trace in memory.
- Dump all logs only on error.

Audit Logs

```
[2014-12-10T21:04:18.251Z] INFO: shakti/restify-audit/21152 on shakti-prod-i-97f6627b: handled: 200 (req_id=16c66e20-80b0-11e4-9f6b-cf625d17b873, audit=true  
, latency=7003, secure=true, _audit=true, req.version=*)  
GET /?locale=en-US&nmrs=1418245445135 HTTP/1.1  
host: www.netflix.com  
x-netflix.appinfo.name: SHAKTI  
accept-language: pt-BR,pt;q=0.8,en-US;q=0.6,en;q=0.4  
x-forwarded-for: 129.170.194.148  
x-forwarded-port: 443  
x-netflix.edge.server.timestamp: 1418245446126  
accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/webp,*/*;q=0.8  
x-netflix.client-host: www.netflix.com  
user-agent: Mozilla/5.0 (Windows NT 6.2; WOW64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/39.0.2171.71 Safari/537.36  
x-forwarded-proto: https  
x-forwarded-host: www.netflix.com  
netflix.nfhttpclient.version: 1.0  
x-netflix-httpclientname: shakti-prod  
  
--  
HTTP/1.1 200 OK  
x-netflix.client.instance: i-97f6627b  
x-frame-options: DENY  
content-type: text/html  
  
--  
req.timers: {  
    "bunyan": 3539,  
    "reqResTracker": 351,  
    "resLocals": 55,  
    "parseQueryString": 4015,  
    "instanceHeaders": 1929,  
    "xForwardedProto": 98,  
    "httpsRedirector": 3353,  
    "readBody": 225,  
    "parseBody": 43,  
    "xframe": 36,  
    "restifyCookieParser": 442,  
    "springbank": 46,  
    "requestContextFp": 3262457,  
    "membershipStatus": 211,  
    "fetchModels": 111865,  
    "htmlContentType": 205,  
    "redirectInAppToEcWeb": 1250,  
    "globalPageRedirect": 51,  
    "openSiteRedirectOrDisplay": 54,  
    "moneyballModeRedirects": 141,  
    "nonMemberContext": 106,  
    "setContextData": 3266,  
    "evidonSetup": 96,  
    "continuumImgUrls": 110,  
    "controlTemplateImgUrls": 190,  
    "ethnioSetup": 6234,  
    "addCspHeader": 43,  
    "redirectCurrentMember": 33,  
    "interimEcWebInterceptor": 270,  
    "render": 3575516  
}
```

URL

status code

request ID (UUID)

req latency

request headers

response headers

individual handler timers

```
req.timers: {
    "bunyan": 3539,
    "reqResTracker": 351,
    "resLocals": 55,
    "parseQueryString": 4015,
    "instanceHeaders": 1929,
    "xForwardedProto": 98,
    "httpsRedirector": 3353,
    "readBody": 225,
    "parseBody": 43,
    "xframe": 36,
    "restifyCookieParser": 442,
    "springbank": 46,
    "requestContextFp": 3262457, ←
    "membershipStatus": 211,
    "fetchModels": 111865,
    "htmlContentType": 205,
    "redirectInAppToEcWeb": 1250,
    "globalPageRedirect": 51,
    "openSiteRedirectOrDisplay": 54,
    "moneyballModeRedirects": 141,
    "nonMemberContext": 106,
    "setContextData": 3266,
    "evidonSetup": 96,
    "continuumImgUrls": 110,
    "controlTemplateImgUrls": 190,
    "ethnioSetup": 6234,
    "addCspHeader": 43,
    "redirectCurrentMember": 33,
    "interimEcWebInterceptor": 270,
    "render": 3575516 ←
```

Scoped Child Loggers

Logging

- Native Bunyan integration.
- Request capture stream.
- Audit logs with each req.
- Scoped child loggers with each req.

restify + Bunyan

- Streaming JSON.
- Processing using Unix tools is easy.
- Helpful tools:
 - Unix: cut(1), wc(1), grep(1), awk(1), perl(1), ...
 - JSON: <https://github.com/trentm/json> (npm install -g json)
 - daggr: <https://github.com/joyent/daggr> (npm install -g daggr)

Examples

- Find all logs for a specific request.
- Count the # of non-200 responses.
- Show all requests that took longer than 200ms.

Advanced Example

Request latency distribution by URL

```
[0] > cat 2015-02-21-shakti-prod-v017-i-232a3529_shakti_shakti.log_0236.log | grep restify-audit| bunyan -0 --strict json -ga latency req.url | awk '{print $1 " " substr($2,0,6)}' | daggr -k 2 -v 1 quantize /WiVie
```

value	Distribution	count
1		0
2		1
4		0
8		0
16		0
32		0
64	@@	43
128	@@@@@@@@@@@	351
256	@@@@@@@@@@@	399
512	@@@	58
1024		5
2048		1
4096		0

/searc

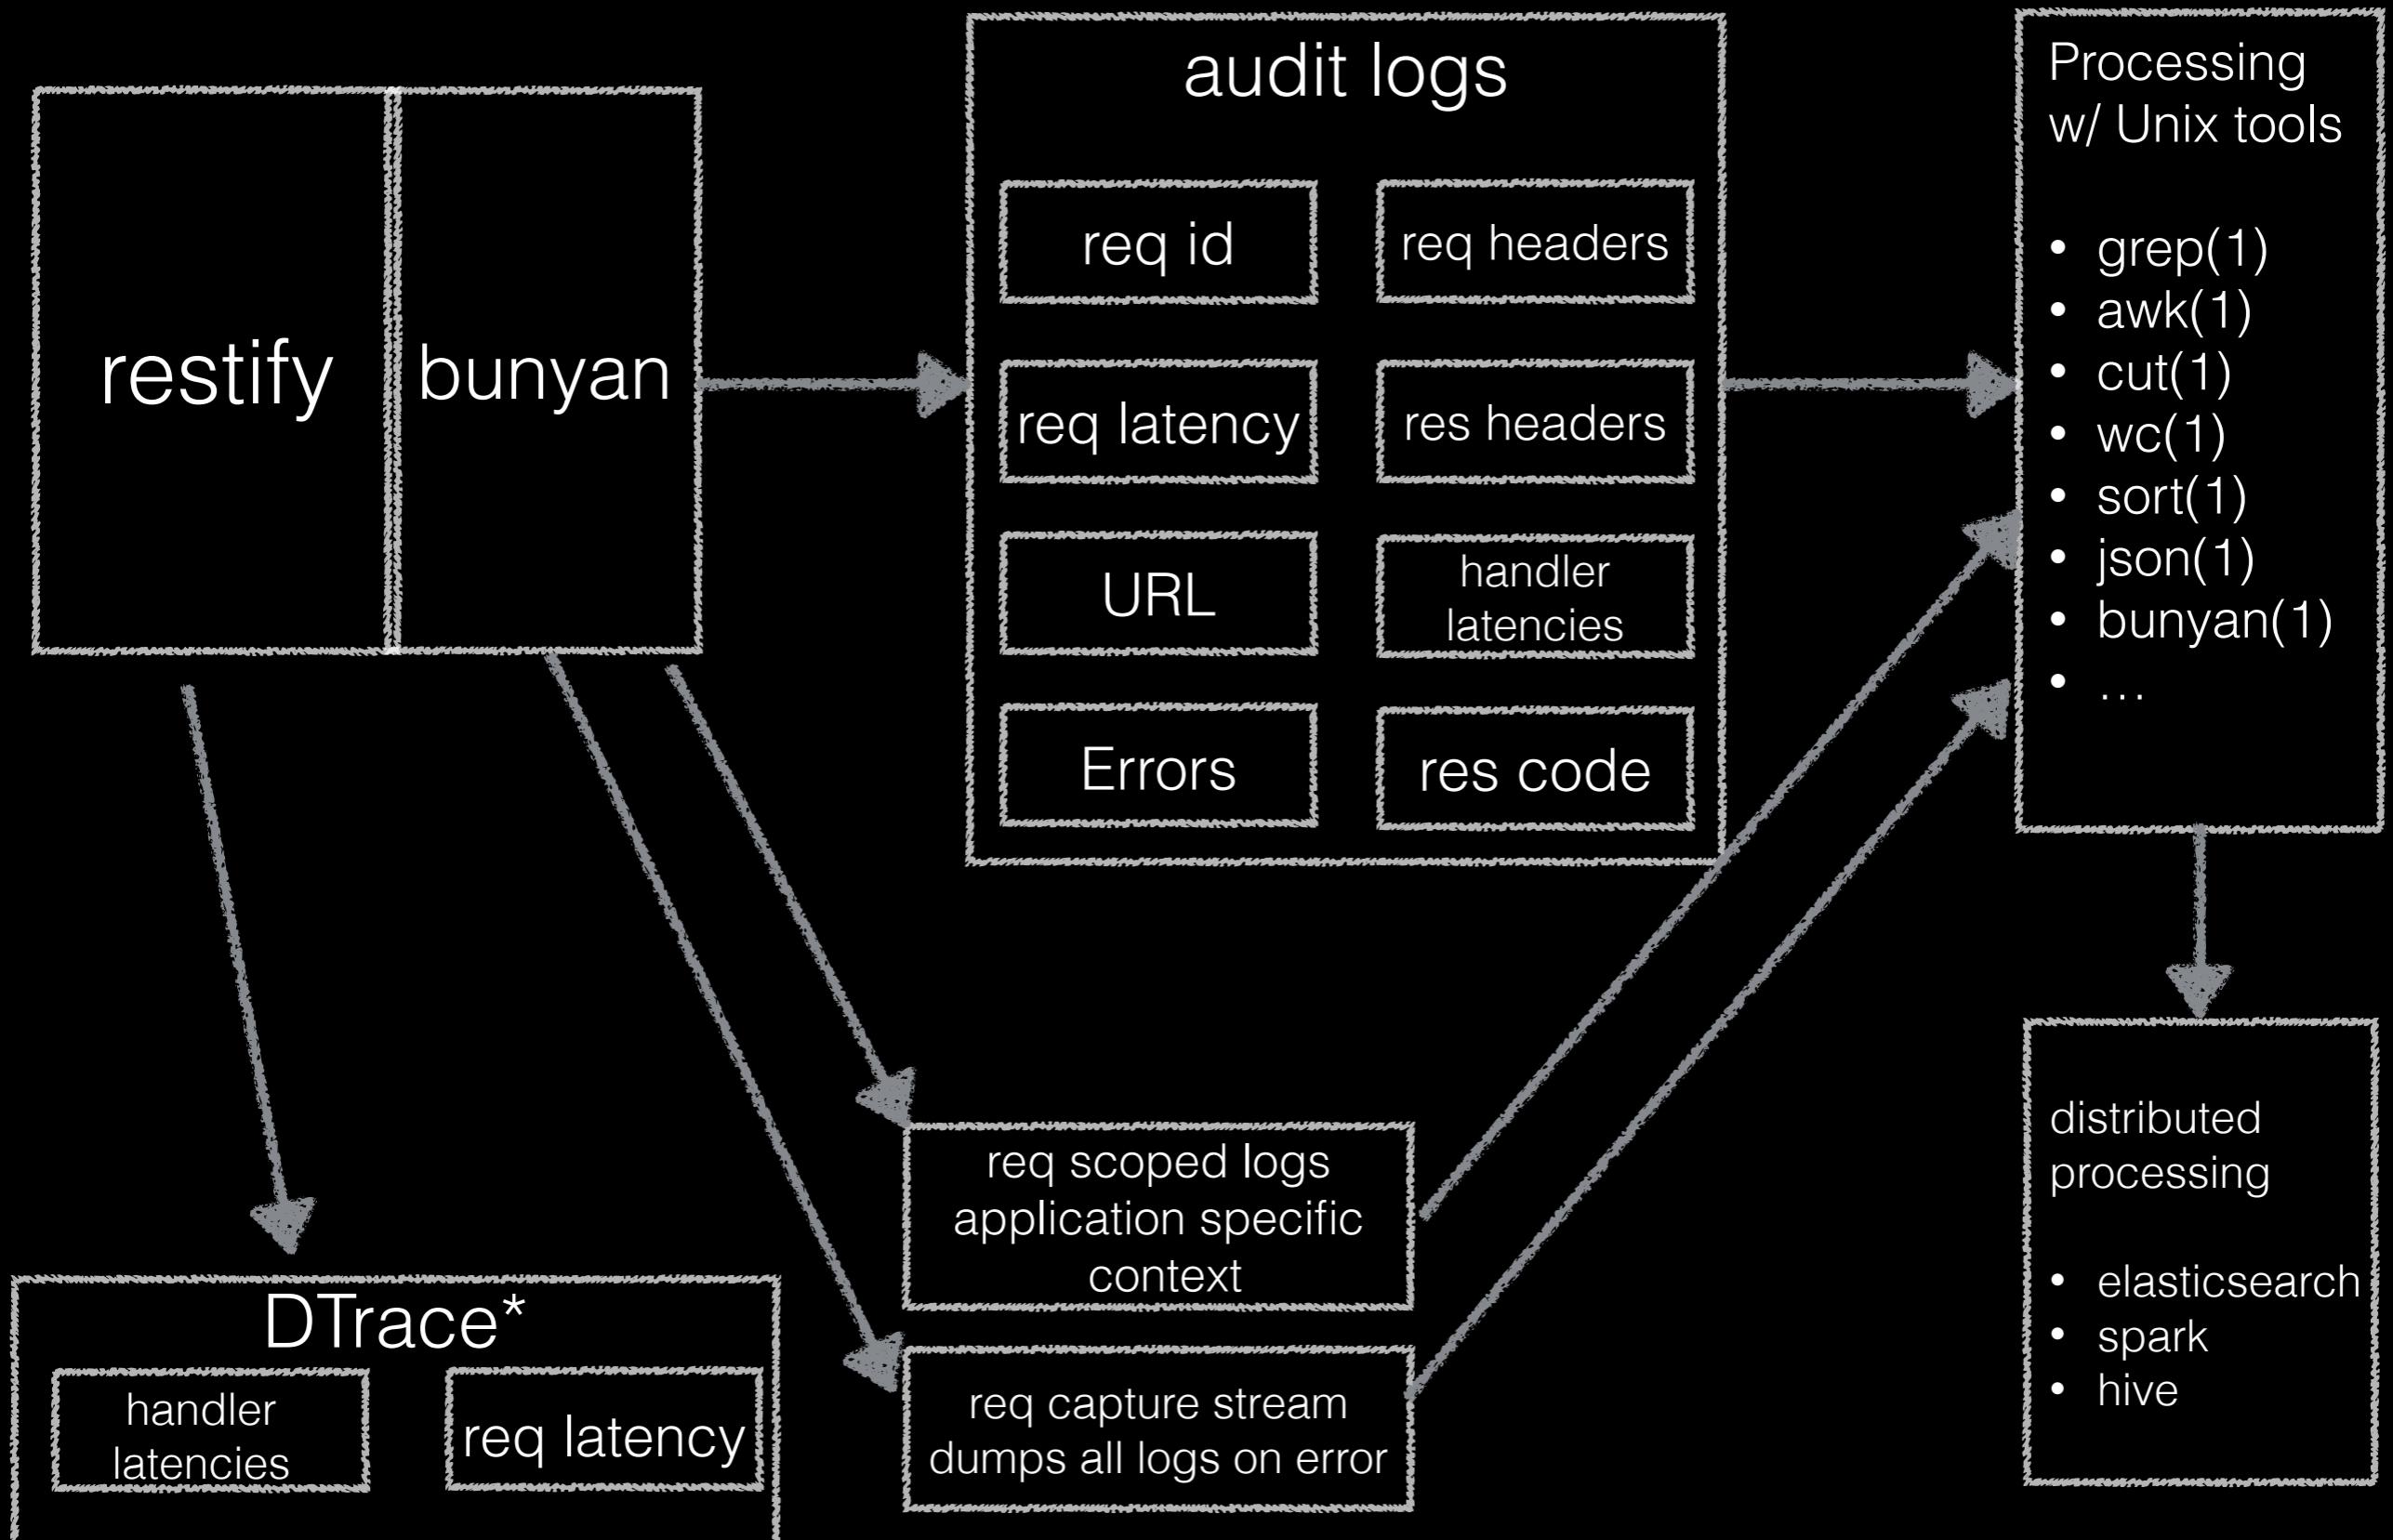
value	Distribution	count
0.5		0
1		13
2		18
4		4
8		1
16		0
32		0
64		0
128	@@@@@@@ @@@@ @@@@ @@@@ @@@@	1285
256	@@@@@@@ @@@@ @@@@ @@@@ @@@@	1270
512	@@	167
1024		15

DTrace

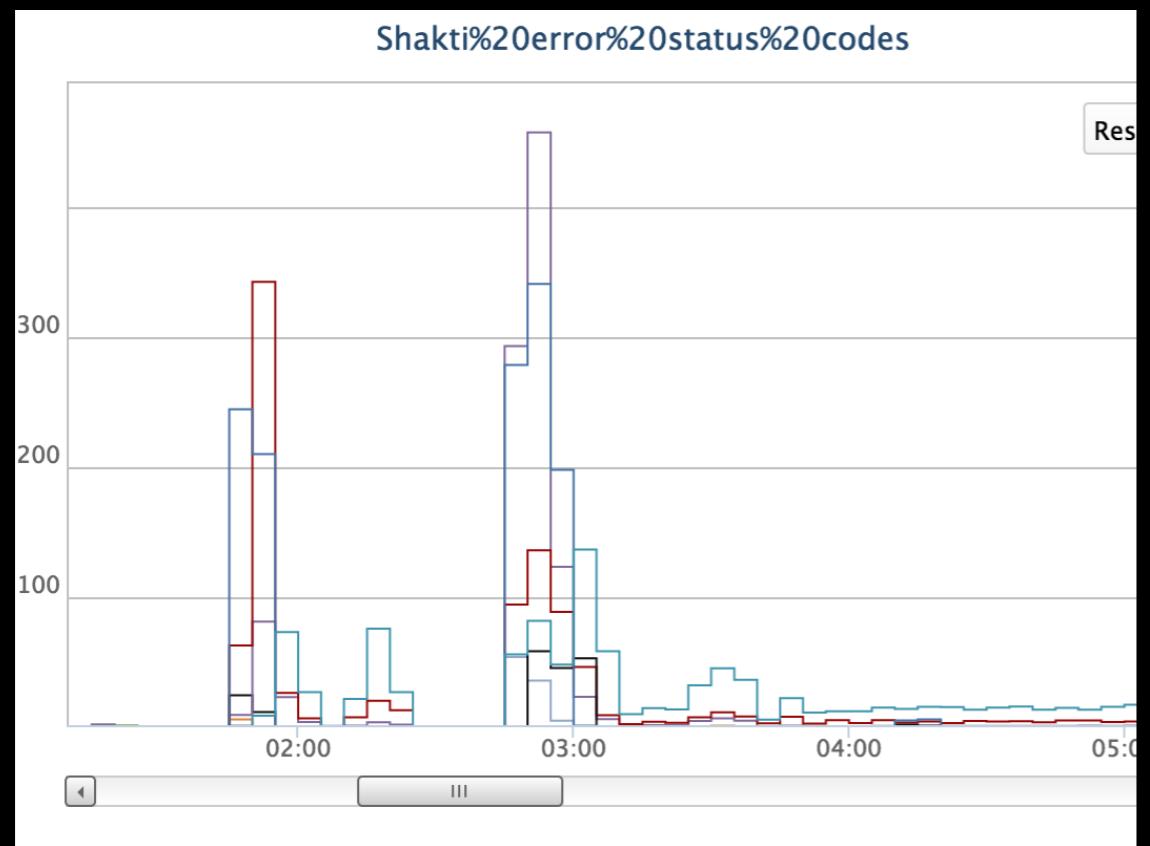
```
└─[yunong@lgml-yunong] - [~] - [2015-03-17 08:15:36]
└─[0] < sudo dtrace -l -P 'restify'
      ID  PROVIDER          MODULE           FUNCTION NAME
12088 restify11084 mod-0x100d05d60 route-start
12089 restify11084 mod-0x100d05d60 handler-start
12090 restify11084 mod-0x100d05d60 handler-done
12091 restify11084 mod-0x100d05d60 route-done
12095 restify11084 mod-0x100d05d60 client-request
16505 restify11084 mod-0x100d05d60 client-response
16506 restify11084 mod-0x100d05d60 client-error
```

Demo

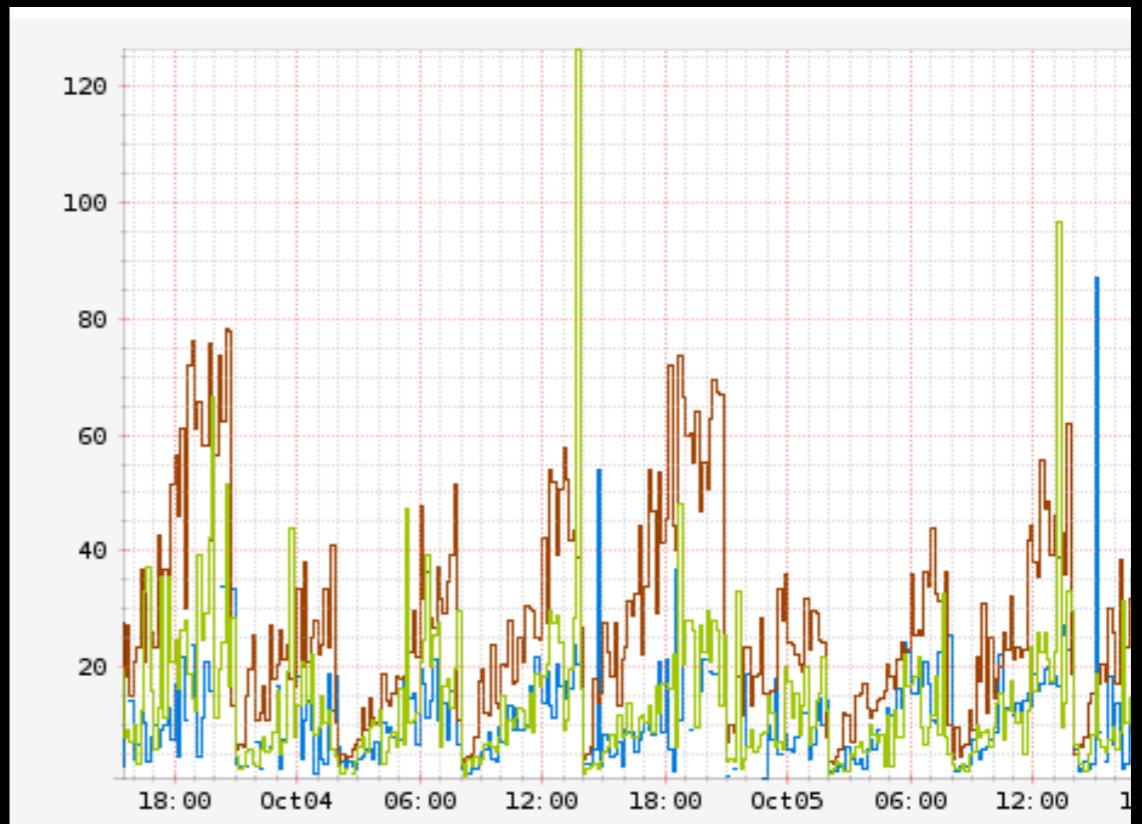
restify metrics



* Where available



Increased Errors



Increased Latency

restify

Interested? Call for contributors.

restifyjs.com

The Killer Combo

- vasync: <https://github.com/davepacheco/node-vasync>
 - Observable async operations.
- bunyan: <https://github.com/trentm/node-bunyan>
 - Streaming JSON logs.
- restify: restifyjs.com
 - Observable REST applications.
- Unix Tools
 - Stream of text model. Many tools to manipulate JSON logs.

Data Data Data





Thanks

- Questions?
- yunong@netflix.com
- [@yunongx](https://twitter.com/yunongx)
- <http://restifyjs.com>

NETFLIX